

**Important:** Please do all projects on `opsys`

## Adding functionality to our system

### Purpose

In the previous project we created a utility that just launched user processes with various limits. In this project we will be adding functionality that we will use in later on projects.

### Task

As in the previous project, you will have two executables, `oss` and `worker`. The `oss` executable will be launching workers. However, we have a few more details.

Our main executable (`oss`) will now be maintaining a "simulated system clock" in shared memory. This system clock is not tied to the main clock of the system, but instead done separately. The clock consists of two separate integers (one storing seconds, the other nanoseconds) in shared memory, both of which are initialized to zero. This system clock must be accessible by the children, so it is required to be in shared memory. The children will not be modifying this clock for this assignment, but they will need to look at it.

In addition to this, `oss` will also maintain a process table (consisting of Process Control Blocks, one for each process). This process table does not need to be in shared memory. The first thing it should keep track of is the PID of the child process, as well as the time right before `oss` does a fork to launch that child process (based on our own simulated clock). It should also contain an entry for if this entry in the process table is empty (ie: not being used). I suggest making your process table an array of structs of PCBs, for example:

```
struct PCB {
    int occupied;           // either true or false
    pid_t pid;              // process id of this child
    int startSeconds;        // time when it was forked
    int startNano;           // time when it was forked
};

struct PCB processTable[20];
```

### worker, the children

The worker takes in two command line arguments, this time corresponding to how many seconds and nanoseconds it should decide to stay around in the system. For example, if you were running it directly you might call it like:

```
./worker 5 500000
```

The worker will then attach to shared memory and examine our simulated system clock. It will then figure out what time it should terminate by adding up the system clock time and the time passed to it. This is when the process should decide to leave the system and terminate.

For example, if the system clock was showing 6 seconds and 100 nanoseconds and the worker was passed 5 and 500000 as above, the target time to terminate in the system would be 11 seconds and 500100 nanoseconds. The worker will then go into a loop, constantly checking the system clock to see if this time has passed. If it ever looks at the system clock and sees values over the ones when it should terminate, it should output some information and then terminate.

So what output should the worker send? Upon starting up, it should output the following information:

```
WORKER PID:6577 PPID:6576 SysClockS: 5 SysclockNano: 1000 TermTimeS: 11 TermTimeNano: 500100
--Just Starting
```

The worker should then go into a loop, checking for its time to expire (IT SHOULD NOT DO A SLEEP). It should also do some periodic output. Everytime it notices that the seconds have changed, it should output a message like:

```
WORKER PID:6577 PPID:6576 SysClockS: 6 SysclockNano: 45000000 TermTimeS: 11 TermTimeNano: 500100
--1 seconds have passed since starting
```

and then one second later it would output:

```
WORKER PID:6577 PPID:6576 SysClockS: 7 SysclockNano: 500000 TermTimeS: 11 TermTimeNano: 500100
--2 seconds have passed since starting
```

Once its time has elapsed, it would send out one final message:

```
WORKER PID:6577 PPID:6576 SysClockS: 11 SysclockNano: 700000 TermTimeS: 11 TermTimeNano: 500100
--Terminating
```

## oss, the parent

The task of oss is to launch a certain number of worker processes with particular parameters. These numbers are determined by its own command line arguments.

Your solution will be invoked using the following command:

```
oss [-h] [-n proc] [-s simul] [-t timelimit]
```

While the first two parameters are similar to the previous project, the -t parameter is different. It now stands for the bound of time that a child process will be launched for. So for example, if it is called with -t 7, then when calling worker processes, it should call them with a time interval randomly between 1 second and 7 seconds (with nanoseconds also random).

When started, oss will initialize the system clock and then go into a loop and start doing a fork() and then an exec() call to launch worker processes. However, it should only do this up to simul number of times. So if called with a -s of 3, we would launch no more than 3 initially. oss should make sure to update the process table with information as it is launching user processes.

This seems close to what we did before, however, will not be doing wait() calls as before. Instead, oss() will be going into a loop, incrementing the clock and then constantly checking to see if a child has terminated. Rough pseudocode for this loop is below:

```
while (stillChildrenToLaunch) {
    incrementClock();

    Every half a second of simulated clock time, output the process table to the screen
```

```

        checkIfChildHasTerminated();

        if (childHasTerminated) {
            updatePCBOfTerminatedChild;
            possiblyLaunchNewChild(obeying process limits)
        }
    }
}

```

The check to see if a child has terminated should be done with a nonblocking `wait()` call. This can be done with code along the lines of:

```
int pid = waitpid(-1, &status, WNOHANG);
```

`waitpid` will return 0 if no child processes have terminated and will return the pid of the child if one has terminated.

The output of `oss` should consist of, every half a second in our simulated system, outputting the entire process table in a nice format. For example:

```

OSS PID:6576 SysClockS: 7 SysclockNano: 500000
Process Table:
Entry Occupied PID  StartS StartN
0      1      6577  5      500000
1      0      0      0      0
2      0      0      0      0
...
19     0      0      0      0

```

## Incrementing the clock

Each iteration in `oss` you need to increment the clock. So how much should you increment it? You should attempt to very loosely have your internal clock be similar to the real clock. This does not have to be precise and does not need to be checked, just use it as a crude guideline. So if you notice that your internal clock is much slower than real time, increase your increment. If it is moving much faster, decrease your increment. Keep in mind that this will change based on server load possibly, so do not worry about if it is off sometimes and on other times.

## Clock race conditions

We are not doing explicit guarding against race conditions. As your child processes are examining the clock, it is possible that the values they see will be slightly off sometimes. This should not be a problem, as the only result would be a clock value that was incorrect very briefly. This could cause the child process to possibly end early, but I will be able to see that as long as your child processes are outputting the clock.

## Signal Handling

In addition, I expect your program to terminate after no more than 60 REAL LIFE seconds. This can be done using a timeout signal, at which point it should kill all currently running child processes and terminate. It should also catch the `ctrl-c` signal, free up shared memory and then terminate all children.

This can be implemented by having your code send a termination signal after 60 seconds. You can find an example of how to do this in our textbook in a source file called `periodicasterik.c`

## Implementation details

It is required for this project that you use version control (git), a `Makefile`, and a `README`. Your `README` file should consist, at a minimum, of a description of how to compile and run your project, any outstanding problems that it still has, and any problems you encountered.

Your makefile should also compile BOTH executables every time. This requires the use of the `all` prefix.

As we are using shared memory, make sure to check and clear out shared memory if you have had errors. Please check the `ipcrm` and `ipcs` commands for this.

## Suggested implementation steps

1. Set up the source files and Makefiles [Day 1]
2. Have oss create shared memory with a clock, then fork off one child and have that child access the shared memory and ensure it works. [Day 2]
3. Write code to parse options and receive the command parameters. [Day 3]
4. Implement the worker taking in an option and doing its loop and output while oss increments. Do not do this in a loop. [Day 4]
5. Implement oss to fork() and then exec() off one worker to do its task and have oss figure out when it is done. [Day 5]
6. Get oss to fork off workers up until the `-n` parameter and do their tasks [Day 6-7]
7. Implement the simultaneous restriction, as well as implement the process table and store data in it. [Day 8-9]
8. Testing and make sure your `README` file indicates how to run your project. Give a one-line example that would let me know how to run it. DO NOT SIMPLY COPY/PASTE THIS DOCUMENT INTO THE README [Day 10+]

## Criteria for success

Please follow the guidelines. Start small, implement one functionality, test. Do not wait until the last minute and contact me if you are having issues.

## Grading

1. *Overall submission: 10 pts.* Program compiles and upon reading, seems to be able to solve the assigned problem.
2. *Code readability: 10 pts.* The code must be readable, with appropriate comments. Author and date should be identified.
3. *Command line parsing: 10 pts.* Program is able to parse the command line appropriately, assigning defaults as needed; issues help if needed.
4. *Makefile: 10 pts.* Makefile works, compiles both files even if both are changed or one is changed. Also ensure your Makefile can do a clean.
5. *README: 10 pts.* Must address any special things you did, or if you missed anything.
6. *Conformance to specifications: 50 pts.* Does your application do the task.

## Submission

Handin an electronic copy of all the sources, `README`, `Makefile(s)`, and results. Create your programs in a directory called `username.1` where `username` is your login name on opsys. Once you are done with everything, *remove the executables and object files*, and issue the following commands:

```
chmod 700 username.2
```

```
cp -p -r username.2 /home/hauschildm/cs4760/assignment2
```

If you have to resubmit, add a `.2` to the end of your directory name and copy that over.

Do not forget `Makefile` (with suffix or pattern rules), your versioning files (`.git` subdirectory), and `README` for the assignment. If you do not use version control, you will lose 10 points. I want to see the log of how the program files are modified. Therefore, you should use some logging mechanism, such as `git`, and let me know about it in your `README`. You must check in the files at least once a day while you are working on them. I do not like to see any extensions on `Makefile` and `README` files.

Before the final submission, perform a `make clean` and keep the latest source checked out in your directory.

You do not have to hand in a hard copy of the project. Assignment is due by 11:59pm on the due date.